

Risc0 Segmentation Test Results

Reza Hasanzadeh

Wholesum Network

rezahsnz@proton.me

July 9th, 2024

1 Introduction

[Risc0](#) is a general-purpose zkVM that enables for arbitrary Rust programs to become verifiable. Given an input Rust program, Risc0 compiles it into a RISC-V binary ELF and feeds into its zk-STARK circuitry producing a verifiable execution trace represented by a *guest* program. This guest program, when run, produces proofs that can be used by third-parties to check if the execution was indeed honest. In order to have infinitely-sized programs, Risc0 divides a guest program into *segments* with segment size capped at 2^{24} *cycles*¹. If a zkVM program does not fit into $\approx 16\text{m}$ cycles, it needs to be partitioned into segments and proved one by one. Segment size can be customized to lower limits too with the acceptable range starting at 2^{13} cycles. Along with execution steps, each segment size has particular memory requirements shown on Table 1 (a comprehensive benchmarking is available [here](#)).

| Segment size <i>cycles</i> | Minimum required memory \approx <i>amount</i> |
|-------------------------------|--|
| $2^{13} \approx 8k$ | 64 MB |
| $2^{14} \approx 16k$ | 128 MB |
| $2^{15} \approx 32k$ | 256 MB |
| $2^{16} \approx 64k$ | 512 MB |
| $2^{17} \approx 128k$ | 1 GB |
| $2^{18} \approx 256k$ | 2 GB |
| $2^{19} \approx 512k$ | 4 GB |
| $2^{20} \approx 1M$ | 8 GB |
| $2^{21} \approx 2M$ | 16 GB |
| $2^{22} \approx 4M$ | 32 GB |
| $2^{23} \approx 8M$ | 64 GB |
| $2^{24} \approx 16M$ | 128 GB |

Table 1: Memory requirements per segment size

Real-world zkVM programs are often big in size. Next section provides a detailed record of inflating and proving a typical guest program. It measures performance and storage requirements of various segment size limits.

¹A [cycle](#) is roughly equivalent to a RISC-V operation.

2 Segmentation test results

In order to emulate a big guest program, a large vector of random u32 numbers were written to the guest memory in each round. While there are more sophisticated inflation methods, this method is easy² and it can be used as a baseline for further experiments. To segment the guest program, we limited its *po2-size-limit* in each round resulting in multiple independent segments per limit size. These segments were then saved as binary blobs to disk for later individual proving. There are 8 rounds of segmentation and proving in total where each round limits the segment size from 2^{15} up to 2^{24} . To measure performance of individual segment proving, a subset of segment blobs were drawn at random(uniform distribution), proved, and the maximum values were recorded. The test computer is a *Lenovo Thinkpad e450(Intel Core i7 5500U(2C 4T), 16GB RAM)* laptop running an *Ubuntu 18.04* dockerized-Risc0-1.0.1 environment.

Rounds

Round 1: up to *1m* cycles

User cycles: *872,543*

Inflation method: write 11k u32 numbers to the guest

The results are displayed on the following table:

| Segment size | 2^{15} | 2^{16} | 2^{17} | 2^{18} | 2^{19} | 2^{20} | 2^{21} | 2^{22} | 2^{23} | 2^{24} |
|-------------------|----------|----------|----------|----------|----------|----------------|----------|----------|----------|----------|
| # Segments | 169 | 26 | 10 | 5 | 3 | 1 | - | - | - | - |
| Total cycles | 5.5M | 1.7M | 1.3M | 1.2M | 1.1M | $\approx 1M^a$ | - | - | - | - |
| Induced inflation | 5.281x | 1.625x | 1.250x | 1.125x | 1.063x | 1.000x | - | - | - | - |
| Prove time | 14s | 28s | 56s | 120s | 262s | 541s | - | - | - | - |
| Sample size | 25% | 50% | 70% | 80% | 100% | 100% | - | - | - | - |
| Blob size | 19kb | 37kb | 52kb | 88kb | 158kb | 298kb | - | - | - | - |

^aExact value: 1,048,576

Table 2: Memory requirements per segment size for a 1M cycles guest

Round 2: up to *2.5m* cycles

User cycles: *1,820,543*

Inflation method: write 23k u32 numbers to the guest

The results are displayed on the following table:

| Segment size | 2^{15} | 2^{16} | 2^{17} | 2^{18} | 2^{19} | 2^{20} | 2^{21} | 2^{22} | 2^{23} | 2^{24} |
|-------------------|----------|----------|----------|----------|----------|----------|------------------|----------|----------|----------|
| # Segments | 333 | 53 | 20 | 9 | 5 | 3 | 1 | - | - | - |
| Total cycles | 11m | 3.5m | 2.6m | 2.4m | 2.2m | 2.2m | $\approx 2.1m^a$ | - | - | - |
| Induced inflation | 5.203x | 1.656x | 1.250x | 1.125x | 1.063x | 1.031x | 1.000x | - | - | - |
| Prove time | 14s | 29s | 58s | 123s | 262s | 540s | - | - | - | - |
| Sample size | 15% | 40% | 50% | 50% | 70% | 100% | - | - | - | - |
| Blob size | 19kb | 35kb | 52kb | 88kb | 158kb | 299kb | 581kb | - | - | - |

^aExact value: 2,097,152

Table 3: Memory requirements per segment size for a 2.5M cycles guest

²RISC architecture tends to have lower number of instructions, so a cycle representing a load instruction, for example, is not inferior to the one representing a mathematical operation.

Round 3: up to $5m$ cycles

User cycles: $3,637,543$

Inflation method: write 46k u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|----------|----------|----------|----------|----------|----------|----------|------------------|----------|----------|
| <i>Segment size</i> | 2^{15} | 2^{16} | 2^{17} | 2^{18} | 2^{19} | 2^{20} | 2^{21} | 2^{22} | 2^{23} | 2^{24} |
| <i># Segments</i> | 647 | 105 | 40 | 18 | 9 | 4 | 2 | 1 | - | - |
| <i>Total cycles</i> | 21.2m | 6.9m | 5.2m | 4.6m | 4.3m | 4.2m | 4.2m | $\approx 4.2m^a$ | - | - |
| <i>Induced inflation</i> | 5.055x | 1.641x | 1.234x | 1.094x | 1.031x | 1.000x | 1.000x | 1.000x | - | - |
| <i>Prove time</i> | 13s | 28s | 57s | 120s | 249s | 536s | - | - | - | - |
| <i>Sample size</i> | 2.5% | 15% | 20% | 25% | 30% | 75% | - | - | - | - |
| <i>Blob size</i> | 19kb | 35kb | 52kb | 88kb | 158kb | 301kb | 581kb | 1.1mb | - | - |

^aExact value: 4,194,304

Table 4: Memory requirements per segment size for a 5M cycles guest

Round 4: up to $10m$ cycles

User cycles: $7,271,543$

Inflation method: write 92k u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|----------|----------|----------|----------|----------|----------|----------|----------|------------------|----------|
| <i>Segment size</i> | 2^{15} | 2^{16} | 2^{17} | 2^{18} | 2^{19} | 2^{20} | 2^{21} | 2^{22} | 2^{23} | 2^{24} |
| <i># Segments</i> | 1,278 | 209 | 78 | 35 | 17 | 8 | 4 | 2 | 1 | - |
| <i>Total cycles</i> | 41.9m | 13.7m | 10.2m | 9.2m | 8.7m | 8.4m | 8.4m | 8.4m | $\approx 8.4m^a$ | - |
| <i>Induced inflation</i> | 4.992x | 1.633x | 1.219x | 1.094x | 1.031x | 1.000x | 1.000x | 1.000x | 1.000x | - |
| <i>Prove time</i> | 14s | 29s | 57s | 119s | 248s | 539s | - | - | - | - |
| <i>Sample size</i> | 3% | 10% | 15% | 15% | 20% | 30% | - | - | - | - |
| <i>Blob size</i> | 19kb | 37kb | 52kb | 88kb | 158kb | 299kb | 581kb | 1.2mb | 2.2mb | - |

^aExact value: 8,388,608

Table 5: Memory requirements per segment size for a 10M cycles guest

Round 5: up to $20m$ cycles

User cycles: $15,013,543$

Inflation method: write 190k u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|----------|----------|----------|----------|----------|----------|----------|----------|----------|-------------------|
| <i>Segment size</i> | 2^{15} | 2^{16} | 2^{17} | 2^{18} | 2^{19} | 2^{20} | 2^{21} | 2^{22} | 2^{23} | 2^{24} |
| <i># Segments</i> | 2,616 | 429 | 161 | 71 | 34 | 17 | 9 | 5 | 3 | 1 |
| <i>Total cycles</i> | 85.7m | 28.1m | 21m | 18.6m | 17.8m | 17.3m | 17m | 16.9m | 16.8m | $\approx 16.8m^a$ |
| <i>Induced inflation</i> | 5.109x | 1.676x | 1.254x | 1.109x | 1.063x | 1.031x | 1.016x | 1.008x | 1.004x | 1.000x |
| <i>Prove time</i> | 14s | 29s | 59s | 121s | 251s | 539s | - | - | - | - |
| <i>Sample size</i> | 1% | 5% | 10% | 10% | 10% | 15% | - | - | - | - |
| <i>Blob size</i> | 19kb | 37kb | 52kb | 88kb | 158kb | 299kb | 581kb | 1.2mb | 2.3mb | 4.5m |

^aExact value: 16,777,216

Table 6: Memory requirements per segment size for a 20M cycles guest

Round 6: up to $45m$ cycles

User cycles: $32,390,116$

Inflation method: write 400k u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-------------------|
| <i>Segment size</i> | 2 ¹⁵ | 2 ¹⁶ | 2 ¹⁷ | 2 ¹⁸ | 2 ¹⁹ | 2 ²⁰ | 2 ²¹ | 2 ²² | 2 ²³ | 2 ²⁴ |
| <i># Segments</i> | 7,187 | 1,019 | 379 | 168 | 80 | 39 | 20 | 10 | 5 | 3 |
| <i>Total cycles</i> | 235.5m | 66.8m | 49.7m | 44m | 41.7m | 40.9m | 40.1m | 39.8m | 41.9m | $\approx 41.9m^a$ |
| <i>Induced inflation</i> | 5.615x | 1.592x | 1.184x | 1.050x | 0.994x | 0.975x | 0.956x | 0.950x | 1.000x | 1.000x |
| <i>Prove time</i> | 15s | 30s | 58s | 119s | 245s | 541s | - | - | - | - |
| <i>Sample size</i> | 0.1% | 2.5% | 5% | 5% | 5% | 10% | - | - | - | - |
| <i>Blob size</i> | 20kb | 36kb | 67kb | 133kb | 265kb | 528kb | 1.1mb | 1.7mb | 2.8mb | 5.2mb |

^aExact value: 41,943,040

Table 7: Memory requirements per segment size for a 45M cycles guest

Round 7: up to 90m cycles

User cycles: 65,563,120

Inflation method: write 800k u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|-----------------|-----------------|-----------------|-----------------|-----------------|-------------------|-----------------|-----------------|-----------------|-------------------|
| <i>Segment size</i> | 2 ¹⁵ | 2 ¹⁶ | 2 ¹⁷ | 2 ¹⁸ | 2 ¹⁹ | 2 ²⁰ | 2 ²¹ | 2 ²² | 2 ²³ | 2 ²⁴ |
| <i># Segments</i> | 16,808 | 2,174 | 801 | 355 | 168 | 82 | 41 | 20 | 10 | 5 |
| <i>Total cycles</i> | 551m | 142.5m | 105m | 92.9m | 87.8m | 85.5m | 84.4m | 83.9m | 83.9m | $\approx 83.9m^a$ |
| <i>Induced inflation</i> | 6.566x | 1.698x | 1.252x | 1.107x | 1.047x | 1.019x | 1.006x | 1.000x | 1.000x | 1.000x |
| <i>Prove time</i> | 14s | 28s | 58s | 119s | 249s | 534s ^b | - | - | - | - |
| <i>Sample size</i> | 0.1% | 2.5% | 5% | 5% | 5% | 3% | - | - | - | - |
| <i>Blob size</i> | 19kb | 38kb | 67kb | 133kb | 265kb | 529kb | 1.1mb | 2mb | 4.1mb | 6.4mb |

^aExact value: 83,886,080

^blow value is due to small sample size

Table 8: Memory requirements per segment size for a 90M cycles guest

Round 8: up to 180m cycles

User cycles: 131,908,988

Inflation method: write 1.6m u32 numbers to the guest

The results are displayed on the following table:

| | | | | | | | | | | |
|--------------------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|------------------|
| <i>Segment size</i> | 2 ¹⁵ | 2 ¹⁶ | 2 ¹⁷ | 2 ¹⁸ | 2 ¹⁹ | 2 ²⁰ | 2 ²¹ | 2 ²² | 2 ²³ | 2 ²⁴ |
| <i># Segments</i> | 36,046 | 4,483 | 1,643 | 727 | 344 | 167 | 83 | 41 | 21 | 11 |
| <i>Total cycles</i> | 1.18b | 293.8m | 215.4m | 190.6m | 180.1m | 175.1m | 173m | 172m | 172m | $\approx 172m^a$ |
| <i>Induced inflation</i> | 6.869x | 1.708x | 1.252x | 1.108x | 1.047x | 1.018x | 1.006x | 1.000x | 1.000x | 1.000x |
| <i>Prove time</i> | 15s | 29s | 66s | 124s | 257s | 544s | - | - | - | - |
| <i>Sample size</i> | 0.1% | 2.5% | 5% | 5% | 5% | 2% | - | - | - | - |
| <i>Blob size</i> | 19kb | 35kb | 67kb | 133kb | 265kb | 529kb | 1.1mb | 2mb | 4.1mb | 7.4mb |

^aExact value: 171,966,464

Table 9: Memory requirements per segment size for a 180M cycles guest

Conclusion

This note investigated a scenario where a Risc0 guest program gets inflated with random inputs. 8 rounds of tests were conducted on segmentation and individual segment proving with a focus on systems whose memory are capped at 8GB. The results point to sweet spots on 2¹⁸ and 2¹⁹ segment size limits where parallelized proving setups could be employed to prove prohibitively large programs quickly with acceptable overhead and inflation in total cycles.